

ETHEREUM FRAUD DETECTION

Problem Statement

- To predict whether the transaction is fraudulent or not-fraudulent using the transaction data.

Dataset Description

- The dataset folder contains the following files:
transaction_dataset.csv : (9841, 51)

Columns Provided in the Dataset

- Index: the index number of a row
- Address: the address of the ethereum account
- FLAG: whether the transaction is fraud or not
- Avg min between sent txn: Average time between sent transactions for account in minutes
- Avgminbetweenreceivedtxn: Average time between received transactions for account in minutes
- TimeDiffbetweenfirstand_last(Mins): Time difference between the first and last transaction
- Sent_txn: Total number of sent normal transactions
- Received_txn: Total number of received normal transactions
- NumberofCreated_Contracts: Total Number of created contract transactions
- UniqueReceivedFrom_Addresses: Total Unique addresses from which account received transactions
- UniqueSentTo_Addresses20: Total Unique addresses from which account sent transactions
- MinValueReceived: Minimum value in Ether ever received
- MaxValueReceived: Maximum value in Ether ever received
- AvgValueReceived5Average value in Ether ever received
- MinValSent: Minimum value of Ether ever sent
- MaxValSent: Maximum value of Ether ever sent
- AvgValSent: Average value of Ether ever sent
- MinValueSentToContract: Minimum value of Ether sent to a contract
- MaxValueSentToContract: Maximum value of Ether sent to a contract
- AvgValueSentToContract: Average value of Ether sent to contracts
- TotalTransactions(IncludingTxnToCreate_Contract): Total number of transactions
- TotalEtherSent: Total Ether sent for account address
- TotalEtherReceived: Total Ether received for account address
- TotalEtherSent_Contracts: Total Ether sent to Contract addresses
- TotalEtherBalance: Total Ether Balance following enacted transactions
- TotalERC20Txns: Total number of ERC20 token transfer transactions
- ERC20TotalEther_Received: Total ERC20 token received transactions in Ether
- ERC20TotalEther_Sent: Total ERC20token sent transactions in Ether
- ERC20TotalEtherSentContract: Total ERC20 token transfer to other contracts in Ether
- ERC20UniqSent_Addr: Number of ERC20 token transactions sent to Unique account addresses
- ERC20UniqRec_Addr: Number of ERC20 token transactions received from Unique addresses
- ERC20UniqRecContractAddr: Number of ERC20token transactions received from Unique contract addresses
- ERC20AvgTimeBetweenSent_Txn: Average time between ERC20 token sent transactions in minutes
- ERC20AvgTimeBetweenRec_Txn: Average time between ERC20 token received transactions in minutes
- ERC20AvgTimeBetweenContract_Txn: Average time ERC20 token between sent token transactions
- ERC20MinVal_Rec: Minimum value in Ether received from ERC20 token transactions for account
- ERC20MaxVal_Rec: Maximum value in Ether received from ERC20 token transactions for account
- ERC20AvgVal_Rec: Average value in Ether received from ERC20 token transactions for account
- ERC20MinVal_Sent: Minimum value in Ether sent from ERC20 token transactions for account
- ERC20MaxVal_Sent: Maximum value in Ether sent from ERC20 token transactions for account
- ERC20AvgVal_Sent: Average value in Ether sent from ERC20 token transactions for account
- ERC20UniqSentTokenName: Number of Unique ERC20 tokens transferred
- ERC20UniqRecTokenName: Number of Unique ERC20 tokens received
- ERC20MostSentTokenType: Most sent token for account via ERC20 transaction
- ERC20MostRecTokenType: Most received token for account via ERC20 transactions

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [2]: # Importing Necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
pd.set_option('display.max_columns', None)
```

```
In [3]: # Read csv file using pandas
path = '/content/drive/MyDrive/Projects/Ethereum Fraud Detection/transaction_dataset.csv'
df = pd.read_csv(path)
```

Basic EDA

```
In [4]: #check shape of the dataset
df.shape
```

```
Out[4]: (9841, 51)
```

```
In [5]: # check columns of dataset
df.columns
```

```
Out[5]: Index(['Unnamed: 0', 'Index', 'Address', 'FLAG', 'Avg min between sent tnx',
              'Avg min between received tnx',
              'Time Diff between first and last (Mins)', 'Sent tnx', 'Received Tnx',
              'Number of Created Contracts', 'Unique Received From Addresses',
              'Unique Sent To Addresses', 'min value received', 'max value received ',
              'avg val received', 'min val sent', 'max val sent', 'avg val sent',
              'min value sent to contract', 'max val sent to contract',
              'avg value sent to contract',
              'total transactions (including tnx to create contract',
              'total Ether sent', 'total ether received',
              'total ether sent contracts', 'total ether balance',
              ' Total ERC20 txns', ' ERC20 total Ether received',
              ' ERC20 total ether sent', ' ERC20 total Ether sent contract',
              ' ERC20 uniq sent addr', ' ERC20 uniq rec addr',
              ' ERC20 uniq sent addr.1', ' ERC20 uniq rec contract addr',
              ' ERC20 avg time between sent tnx', ' ERC20 avg time between rec tnx',
              ' ERC20 avg time between rec 2 tnx',
              ' ERC20 avg time between contract tnx', ' ERC20 min val rec',
              ' ERC20 max val rec', ' ERC20 avg val rec', ' ERC20 min val sent',
              ' ERC20 max val sent', ' ERC20 avg val sent',
              ' ERC20 min val sent contract', ' ERC20 max val sent contract',
              ' ERC20 avg val sent contract', ' ERC20 uniq sent token name',
              ' ERC20 uniq rec token name', ' ERC20 most sent token type',
              ' ERC20_most_rec_token_type'],
              dtype='object')
```

```
In [6]: #drop Unnamed: 0, Address, Index from dataset
df.drop(columns = ['Unnamed: 0', 'Address', 'Index'], inplace = True)
```

```
In [7]: # Check which columns are having categorical, numerical or boolean values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9841 entries, 0 to 9840
Data columns (total 48 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   FLAG                                                                    9841 non-null   int64
1   Avg min between sent tnx                                                9841 non-null   float64
2   Avg min between received tnx                                            9841 non-null   float64
3   Time Diff between first and last (Mins)                                9841 non-null   float64
4   Sent tnx                                                                9841 non-null   int64
5   Received Tnx                                                            9841 non-null   int64
6   Number of Created Contracts                                             9841 non-null   int64
7   Unique Received From Addresses                                         9841 non-null   int64
8   Unique Sent To Addresses                                               9841 non-null   int64
9   min value received                                                      9841 non-null   float64
10  max value received                                                      9841 non-null   float64
11  avg val received                                                        9841 non-null   float64
12  min val sent                                                            9841 non-null   float64
13  max val sent                                                            9841 non-null   float64
14  avg val sent                                                            9841 non-null   float64
15  min value sent to contract                                              9841 non-null   float64
16  max val sent to contract                                                9841 non-null   float64
17  avg value sent to contract                                              9841 non-null   float64
18  total transactions (including tnx to create contract                    9841 non-null   int64
19  total Ether sent                                                        9841 non-null   float64
20  total ether received                                                    9841 non-null   float64
21  total ether sent contracts                                              9841 non-null   float64
22  total ether balance                                                     9841 non-null   float64
23  Total ERC20 txns                                                        9012 non-null   float64
24  ERC20 total Ether received                                              9012 non-null   float64
25  ERC20 total ether sent                                                  9012 non-null   float64
26  ERC20 total Ether sent contract                                         9012 non-null   float64
27  ERC20 uniq sent addr                                                    9012 non-null   float64
28  ERC20 uniq rec addr                                                     9012 non-null   float64
29  ERC20 uniq sent addr.1                                                  9012 non-null   float64
30  ERC20 uniq rec contract addr                                             9012 non-null   float64
31  ERC20 avg time between sent tnx                                         9012 non-null   float64
32  ERC20 avg time between rec tnx                                          9012 non-null   float64
33  ERC20 avg time between rec 2 tnx                                        9012 non-null   float64
34  ERC20 avg time between contract tnx                                     9012 non-null   float64
35  ERC20 min val rec                                                       9012 non-null   float64
36  ERC20 max val rec                                                       9012 non-null   float64
37  ERC20 avg val rec                                                       9012 non-null   float64
38  ERC20 min val sent                                                      9012 non-null   float64
39  ERC20 max val sent                                                      9012 non-null   float64
40  ERC20 avg val sent                                                      9012 non-null   float64
41  ERC20 min val sent contract                                              9012 non-null   float64
42  ERC20 max val sent contract                                              9012 non-null   float64
43  ERC20 avg val sent contract                                              9012 non-null   float64
44  ERC20 uniq sent token name                                              9012 non-null   float64
45  ERC20 uniq rec token name                                               9012 non-null   float64
46  ERC20 most sent token type                                              9000 non-null   object
47  ERC20_most_rec_token_type                                               8990 non-null   object
dtypes: float64(39), int64(7), object(2)
memory usage: 3.6+ MB
```

```
In [8]: # For more information on the dataset Like the total count in all the columns of train data
# min, max values and more information of the respective columns
df.describe().T
```

Out[8]:

	count	mean	std	min	25%	50%	75%	max
FLAG	9841.0	2.214206e-01	4.152241e-01	0.00	0.000000	0.000000e+00	0.000000	1.000000e+00
Avg min between sent txn	9841.0	5.086879e+03	2.148655e+04	0.00	0.000000	1.734000e+01	565.470000	4.302877e+05
Avg min between received txn	9841.0	8.004851e+03	2.308171e+04	0.00	0.000000	5.097700e+02	5480.390000	4.821755e+05
Time Diff between first and last (Mins)	9841.0	2.183333e+05	3.229379e+05	0.00	316.930000	4.663703e+04	304070.980000	1.954861e+06
Sent txn	9841.0	1.159317e+02	7.572264e+02	0.00	1.000000	3.000000e+00	11.000000	1.000000e+04
Received Txn	9841.0	1.637009e+02	9.408366e+02	0.00	1.000000	4.000000e+00	27.000000	1.000000e+04
Number of Created Contracts	9841.0	3.729702e+00	1.414456e+02	0.00	0.000000	0.000000e+00	0.000000	9.995000e+03
Unique Received From Addresses	9841.0	3.036094e+01	2.986211e+02	0.00	1.000000	2.000000e+00	5.000000	9.999000e+03
Unique Sent To Addresses	9841.0	2.584016e+01	2.638204e+02	0.00	1.000000	2.000000e+00	3.000000	9.287000e+03
min value received	9841.0	4.384515e+01	3.259291e+02	0.00	0.001000	9.585600e-02	2.000000	1.000000e+04
max value received	9841.0	5.231525e+02	1.300882e+04	0.00	1.000000	6.000000e+00	67.067040	8.000000e+05
avg val received	9841.0	1.007117e+02	2.885002e+03	0.00	0.426905	1.729730e+00	22.000000	2.836188e+05
min val sent	9841.0	4.800090e+00	1.386097e+02	0.00	0.000000	4.912600e-02	0.998800	1.200000e+04
max val sent	9841.0	3.146173e+02	6.629213e+03	0.00	0.164577	4.999380e+00	61.520653	5.200000e+05
avg val sent	9841.0	4.475573e+01	2.390802e+02	0.00	0.086184	1.606000e+00	21.999380	1.200000e+04
min value sent to contract	9841.0	3.048471e-06	2.253968e-04	0.00	0.000000	0.000000e+00	0.000000	2.000000e-02
max val sent to contract	9841.0	7.725739e-06	5.158151e-04	0.00	0.000000	0.000000e+00	0.000000	4.602900e-02
avg value sent to contract	9841.0	5.387054e-06	3.234341e-04	0.00	0.000000	0.000000e+00	0.000000	2.301400e-02
total transactions (including txn to create contract	9841.0	2.833624e+02	1.352404e+03	0.00	4.000000	8.000000e+00	54.000000	1.999500e+04
total Ether sent	9841.0	1.016092e+04	3.583227e+05	0.00	0.226206	1.248680e+01	100.998974	2.858096e+07
total ether received	9841.0	1.163832e+04	3.642048e+05	0.00	2.670424	3.052963e+01	101.000000	2.858159e+07
total ether sent contracts	9841.0	7.725710e-06	5.158125e-04	0.00	0.000000	0.000000e+00	0.000000	4.602871e-02
total ether balance	9841.0	1.477395e+03	2.424254e+05	-15605352.04	0.000621	1.722000e-03	0.044520	1.428864e+07
Total ERC20 txns	9012.0	3.625566e+01	4.475289e+02	0.00	0.000000	1.000000e+00	3.000000	1.000100e+04
ERC20 total Ether received	9012.0	1.296207e+08	1.053858e+10	0.00	0.000000	1.000000e-12	100.337000	1.000020e+12
ERC20 total ether sent	9012.0	1.386849e+07	1.180390e+09	0.00	0.000000	0.000000e+00	0.000000	1.120000e+11
ERC20 total Ether sent contract	9012.0	1.109392e+02	6.128635e+03	0.00	0.000000	0.000000e+00	0.000000	4.160000e+05
ERC20 uniq sent addr	9012.0	5.638038e+00	1.052525e+02	0.00	0.000000	0.000000e+00	0.000000	6.582000e+03
ERC20 uniq rec addr	9012.0	7.598535e+00	8.181847e+01	0.00	0.000000	1.000000e+00	2.000000	4.293000e+03
ERC20 uniq sent addr.1	9012.0	3.439858e-03	6.569787e-02	0.00	0.000000	0.000000e+00	0.000000	3.000000e+00
ERC20 uniq rec contract addr	9012.0	4.901909e+00	1.724658e+01	0.00	0.000000	1.000000e+00	2.000000	7.820000e+02
ERC20 avg time between sent txn	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 avg time between rec txn	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 avg time between rec 2 txn	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 avg time between contract txn	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 min val rec	9012.0	4.856147e+02	1.688328e+04	0.00	0.000000	0.000000e+00	0.001523	9.900000e+05
ERC20 max val rec	9012.0	1.252524e+08	1.053741e+10	0.00	0.000000	0.000000e+00	99.000000	1.000000e+12
ERC20 avg val rec	9012.0	4.346203e+06	2.141192e+08	0.00	0.000000	0.000000e+00	29.464673	1.724181e+10
ERC20 min val sent	9012.0	1.174126e+04	1.053567e+06	0.00	0.000000	0.000000e+00	0.000000	1.000000e+08
ERC20 max val sent	9012.0	1.303594e+07	1.179905e+09	0.00	0.000000	0.000000e+00	0.000000	1.120000e+11
ERC20 avg val sent	9012.0	6.318389e+06	5.914764e+08	0.00	0.000000	0.000000e+00	0.000000	5.614756e+10
ERC20 min val sent contract	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 max val sent contract	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 avg val sent contract	9012.0	0.000000e+00	0.000000e+00	0.00	0.000000	0.000000e+00	0.000000	0.000000e+00
ERC20 uniq sent token name	9012.0	1.384931e+00	6.735121e+00	0.00	0.000000	0.000000e+00	0.000000	2.130000e+02
ERC20 uniq rec token name	9012.0	4.826676e+00	1.667861e+01	0.00	0.000000	1.000000e+00	2.000000	7.370000e+02

```
In [9]: # check Length of dataset
print(f" dataset length : {len(df)}")

dataset length : 9841
```

```
In [10]: # Loop through dataset to find count of unique values of each column
for col in df.columns:
    print(f"{col} : {len(df[col].value_counts())}")
```

```
FLAG : 2
Avg min between sent tnx : 5013
Avg min between received tnx : 6223
Time Diff between first and last (Mins) : 7810
Sent tnx : 641
Received Tnx : 727
Number of Created Contracts : 20
Unique Received From Addresses : 256
Unique Sent To Addresses : 258
min value received : 4589
max value received : 6302
avg val received : 6767
min val sent : 4719
max val sent : 6647
avg val sent : 5854
min value sent to contract : 3
max val sent to contract : 4
avg value sent to contract : 4
total transactions (including tnx to create contract : 897
total Ether sent : 5868
total ether received : 6728
total ether sent contracts : 4
total ether balance : 5717
Total ERC20 txns : 300
ERC20 total Ether received : 3460
ERC20 total ether sent : 1415
ERC20 total Ether sent contract : 29
ERC20 uniq sent addr : 107
ERC20 uniq rec addr : 147
ERC20 uniq sent addr.1 : 4
ERC20 uniq rec contract addr : 123
ERC20 avg time between sent tnx : 1
ERC20 avg time between rec tnx : 1
ERC20 avg time between rec 2 tnx : 1
ERC20 avg time between contract tnx : 1
ERC20 min val rec : 1276
ERC20 max val rec : 2647
ERC20 avg val rec : 3380
ERC20 min val sent : 476
ERC20 max val sent : 1130
ERC20 avg val sent : 1309
ERC20 min val sent contract : 1
ERC20 max val sent contract : 1
ERC20 avg val sent contract : 1
ERC20 uniq sent token name : 70
ERC20 uniq rec token name : 121
ERC20 most sent token type : 305
ERC20_most_rec_token_type : 467
```

```
In [11]: # Check for missing values in all the columns of the dataset
df.isnull().sum()
```

```
Out[11]: FLAG 0
Avg min between sent tnx 0
Avg min between received tnx 0
Time Diff between first and last (Mins) 0
Sent tnx 0
Received Tnx 0
Number of Created Contracts 0
Unique Received From Addresses 0
Unique Sent To Addresses 0
min value received 0
max value received 0
avg val received 0
min val sent 0
max val sent 0
avg val sent 0
min value sent to contract 0
max val sent to contract 0
avg value sent to contract 0
total transactions (including tnx to create contract 0
total Ether sent 0
total ether received 0
total ether sent contracts 0
total ether balance 0
Total ERC20 txns 829
ERC20 total Ether received 829
ERC20 total ether sent 829
ERC20 total Ether sent contract 829
ERC20 uniq sent addr 829
ERC20 uniq rec addr 829
ERC20 uniq sent addr.1 829
ERC20 uniq rec contract addr 829
ERC20 avg time between sent tnx 829
ERC20 avg time between rec tnx 829
ERC20 avg time between rec 2 tnx 829
ERC20 avg time between contract tnx 829
ERC20 min val rec 829
ERC20 max val rec 829
ERC20 avg val rec 829
ERC20 min val sent 829
ERC20 max val sent 829
ERC20 avg val sent 829
ERC20 min val sent contract 829
ERC20 max val sent contract 829
ERC20 avg val sent contract 829
ERC20 uniq sent token name 829
ERC20 uniq rec token name 829
ERC20 most sent token type 841
ERC20_most_rec_token_type 851
dtype: int64
```

Correlation Matrix

Why?

- A correlation matrix is a table showing correlation coefficients between variables

There are three broad reasons for computing a correlation matrix:

1. To summarize a large amount of data where the goal is to see patterns. In our example above, the observable pattern is that all the variables highly correlate with each other.
2. To input into other analyses. For example, people commonly use correlation matrixes as inputs for exploratory factor analysis, confirmatory factor analysis, structural equation models, and linear regression when excluding missing values pairwise.
3. As a diagnostic when checking other analyses. For example, with linear regression, a high amount of correlations suggests that the linear regression estimates will be unreliable.

```
In [12]: # Using pandas  
df.corr().style.background_gradient(cmap='coolwarm')
```

```
/usr/local/lib/python3.8/dist-packages/pandas/io/formats/style.py:2813: RuntimeWarning: All-NaN slice encountered
  smin = np.nanmin(gmap) if vmin is None else vmin
/usr/local/lib/python3.8/dist-packages/pandas/io/formats/style.py:2814: RuntimeWarning: All-NaN slice encountered
  smax = np.nanmax(gmap) if vmax is None else vmax
```


Out[12]:

	FLAG	Avg min between sent tnx	Avg min between received tnx	Time Diff between first and last (Mins)	Sent tnx	Received Tnx	Number of Created Contracts	Unique Received From Addresses	Unique Sent To Addresses	min value received	max value received	avg val received	min val sent	max val sent	avg val sent	min value sent to contract	max val sent to contract	avg value sent to contract	total transactions (including tnx to create contract
FLAG	1.000000	-0.029754	-0.118533	-0.269354	-0.078006	-0.079316	-0.013711	-0.031941	-0.045584	-0.021641	-0.019259	-0.011881	0.006626	-0.022437	-0.063556	-0.007213	-0.007988	-0.008883	-0.100289
Avg min between sent tnx	-0.029754	1.000000	0.060979	0.214722	-0.032289	-0.035735	-0.006186	-0.015912	-0.017688	-0.014886	-0.007104	-0.004382	-0.000789	-0.005716	0.003597	-0.000210	0.001044	0.000759	-0.043586
Avg min between received tnx	-0.118533	0.060979	1.000000	0.303897	-0.040419	-0.053478	-0.008378	-0.029571	-0.025747	-0.045753	-0.011575	-0.009313	-0.008761	-0.012176	-0.040011	-0.003916	-0.003230	-0.003940	-0.060711
Time Diff between first and last (Mins)	-0.269354	0.214722	0.303897	1.000000	0.154480	0.148376	-0.003881	0.037043	0.071140	-0.084996	-0.002240	-0.014002	-0.013107	0.014194	-0.046039	0.001734	0.036162	0.029440	0.189311
Sent tnx	-0.078006	-0.032289	-0.040419	0.154480	1.000000	0.198455	0.320603	0.130064	0.670014	0.024015	0.102109	0.140677	-0.004846	0.225356	0.027468	0.029529	0.013263	0.020865	0.731503
Received Tnx	-0.079316	-0.035735	-0.053478	0.148376	0.198455	1.000000	-0.003838	0.648655	0.164112	-0.022936	0.224805	-0.001786	0.093448	0.097769	0.125075	0.020645	0.008243	0.013767	0.806393
Number of Created Contracts	-0.013711	-0.006186	-0.008378	-0.003881	0.320603	-0.003838	1.000000	-0.002211	0.084598	-0.001542	-0.000752	-0.000498	-0.000913	0.141919	-0.001083	-0.000261	-0.000225	-0.000270	0.281428
Unique Received From Addresses	-0.031941	-0.015912	-0.029571	0.037043	0.130064	0.648655	-0.002211	1.000000	0.159829	-0.012939	0.175860	-0.001324	0.296240	0.058060	0.226712	0.053946	0.023258	0.037343	0.523848
Unique Sent To Addresses	-0.045584	-0.017688	-0.025747	0.071140	0.670014	0.164112	0.084598	0.159829	1.000000	0.070145	0.148182	0.207410	-0.003166	0.196573	0.022143	0.023183	0.010926	0.016790	0.498165
min value received	-0.021641	-0.014886	-0.045753	-0.084996	0.024015	-0.022936	-0.001542	-0.012939	0.070145	1.000000	0.026710	0.122911	0.117682	0.015061	0.267280	-0.001820	-0.002015	-0.002241	-0.002671
max value received	-0.019259	-0.007104	-0.011575	-0.002240	0.102109	0.224805	-0.000752	0.175860	0.148182	0.026710	1.000000	0.622959	0.002239	0.135937	0.043989	-0.000427	-0.000550	-0.000588	0.213485
avg val received	-0.011881	-0.004382	-0.009313	-0.014002	0.140677	-0.001786	-0.000498	-0.001324	0.207410	0.122911	0.622959	1.000000	0.012988	0.134113	0.070700	-0.000457	-0.000516	-0.000571	0.077472
min val sent	0.006626	-0.000789	-0.008761	-0.013107	-0.004846	0.093448	-0.000913	0.296240	-0.003166	0.117682	0.002239	0.012988	1.000000	0.022662	0.594868	-0.000468	-0.000519	-0.000577	0.062200
max val sent	-0.022437	-0.005716	-0.012176	0.014194	0.225356	0.097769	0.141919	0.058060	0.196573	0.015061	0.135937	0.134113	0.022662	1.000000	0.184962	-0.000378	-0.000594	-0.000605	0.209038
avg val sent	-0.063556	0.003597	-0.040011	-0.046039	0.027468	0.125075	-0.001083	0.226712	0.022143	0.267280	0.043989	0.070700	0.594868	0.184962	1.000000	-0.002445	-0.002764	-0.003056	0.102278
min value sent to contract	-0.007213	-0.000210	-0.003916	0.001734	0.029529	0.020645	-0.000261	0.053946	0.023183	-0.001820	-0.000427	-0.000457	-0.000468	-0.000378	-0.002445	1.000000	0.436849	0.696789	0.030868
max val sent to contract	-0.007988	0.001044	-0.003230	0.036162	0.013263	0.008243	-0.000225	0.023258	0.010926	-0.002015	-0.000550	-0.000516	-0.000519	-0.000594	-0.002764	0.436849	1.000000	0.949607	0.013137
avg value sent to contract	-0.008883	0.000759	-0.003940	0.029440	0.020865	0.013767	-0.000270	0.037343	0.016790	-0.002241	-0.000588	-0.000571	-0.000577	-0.000605	-0.003056	0.696789	0.949607	1.000000	0.021232
total transactions (including tnx to create contract	-0.100289	-0.043586	-0.060711	0.189311	0.731503	0.806393	0.281428	0.523848	0.498165	-0.002671	0.213485	0.077472	0.062200	0.209038	0.102278	0.030868	0.013137	0.021232	1.000000
total Ether sent	-0.014993	-0.006440	-0.009442	0.012999	0.244434	0.132150	0.024877	0.028881	0.164535	-0.001748	0.112739	0.155433	0.000166	0.470139	0.198750	-0.000226	-0.000356	-0.000363	0.231397
total ether received	-0.016900	-0.007285	-0.010720	0.014756	0.155811	0.235282	-0.000753	0.124897	0.086376	-0.001874	0.296173	0.062609	0.001171	0.086551	0.165232	-0.000252	-0.000400	-0.000406	0.250842
total ether sent contracts	-0.007988	0.001044	-0.003230	0.036162	0.013263	0.008243	-0.000225	0.023258	0.010926	-0.002015	-0.000550	-0.000516	-0.000519	-0.000594	-0.002764	0.436852	1.000000	0.949608	0.013137
total ether balance	-0.003229	-0.001425	-0.002149	0.002955	-0.127211	0.158146	-0.037902	0.144949	-0.113428	-0.000233	0.278315	-0.135682	0.001514	-0.564872	-0.045534	-0.000043	-0.000074	-0.000074	0.034828
Total ERC20 tnxs	-0.025697	-0.012307	-0.020578	0.075922	0.381311	0.116203	0.249500	0.057718	0.148071	-0.007648	0.032261	0.012264	-0.002574	0.130834	0.020393	-0.000103	-0.000304	-0.000279	0.320750
ERC20 total Ether received	-0.004475	-0.002542	-0.003656	0.046788	0.006683	0.016501	0.002072	0.002742	0.005709	-0.001727	0.006230	0.004474	-0.000425	0.032880	0.019350	0.001242	0.000426	0.000773	0.015447
ERC20 total ether sent	0.024762	-0.002105	-0.003876	-0.005014	-0.000315	-0.000452	0.001166	0.000059	0.000004	-0.001627	-0.000266	-0.000390	-0.000407	0.000615	-0.001598	-0.000157	-0.000180	-0.000198	-0.000369
ERC20 total Ether sent contract	0.013514	-0.004409	-0.006160	-0.008389	-0.002897	0.104870	-0.000374	0.116672	-0.001848	-0.002546	0.068704	-0.000330	-0.000628	-0.000896	-0.003523	-0.000256	-0.000283	-0.000315	0.071317
ERC20 uniq sent addr	-0.020554	-0.009477	-0.014776	0.040686	0.356065	0.044187	0.154788	0.045200	0.121789	-0.006072	0.001475	-0.000135	-0.001857	0.046650	-0.006726	-0.000578	0.000050	-0.000162	0.246555
ERC20 uniq rec addr	-0.017539	-0.001133	-0.016037	0.080709	0.300461	0.136485	0.190646	0.083026	0.176988	-0.005190	0.039185	0.037338	-0.002092	0.198918	0.010298	0.002891	0.001310	0.002052	0.283377
ERC20 uniq sent addr.1	-0.003047	-0.012754	-0.014698	0.019961	-0.008379	0.337727	-0.001083	0.241494	-0.005345	-0.007365	0.048690	-0.001739	-0.001817	-0.002590	-0.010191	-0.000740	-0.000820	-0.000911	0.230217
ERC20 uniq rec contract addr	-0.018527	0.045615	-0.014932	0.319176	0.219665	0.201686	0.028986	0.147741	0.237351	-0.003216	0.181375	0.202583	-0.004108	0.165755	0.045245	0.018657	0.012053	0.016112	0.266519
ERC20 avg time between sent tnx	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 avg time between rec tnx	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 avg time between rec 2 tnx	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 avg time between contract tnx	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 min val rec	0.011163	0.004673	-0.008319	-0.010543	-0.003868	-0.004464	-0.000788	-0.001999	-0.002333	-0.001154	-0.001076	-0.000648	-0.000490	-0.001124	-0.002931	-0.000406	-0.000450	-0.000500	-0.005357
ERC20 max val rec	-0.004313	-0.002456	-0.003544	0.046512	0.004291	0.015379	0.000920	0.001603	0.003598	-0.001670	0.003993	0.001357	-0.000411	0.031620	0.018990	0.000526	0.000117	0.000277	0.013204
ERC20 avg val rec	0.007881	-0.003149	-0.005606	0.049020	0.008843	0.019126	0.002052	0.003592	0.005246	-0.002852	0.002440	-0.000120	-0.000694	0.026005	0.014155	0.000573	0.000058	0.000246	0.018482
ERC20 min val sent	0.025364	-0.001687	-0.003758	-0.006852	-0.001557	-0.001899	-0.000092	-0.001044	-0.000967	-0.001498	-0.000448	-0.000386	-0.000387	-0.000514	-0.001974	-0.000157	-0.000174	-0.000194	-0.002204

	FLAG	Avg min between sent tnx	Avg min between received tnx	Time Diff between first and last (Mins)	Sent tnx	Received Tnx	Number of Created Contracts	Unique Received From Addresses	Unique Sent To Addresses	min value received	max value received	avg val received	min val sent	max val sent	avg val sent	min value sent to contract	max val sent to contract	avg value sent to contract	total transactions (including tnx to create contract
ERC20 max val sent	0.025038	-0.002020	-0.003780	-0.006268	-0.001019	-0.000744	0.000149	-0.000261	-0.000439	-0.001529	-0.000293	-0.000380	-0.000383	-0.000194	-0.001682	-0.000148	-0.000169	-0.000186	-0.001073
ERC20 avg val sent	0.025044	-0.001965	-0.003727	-0.006802	-0.001415	-0.000918	-0.000070	-0.000546	-0.000827	-0.001477	-0.000315	-0.000377	-0.000370	-0.000413	-0.001667	-0.000147	-0.000165	-0.000183	-0.001439
ERC20 min val sent contract	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 max val sent contract	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 avg val sent contract	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan	nan
ERC20 uniq sent token name	0.009249	0.000529	-0.019896	0.263939	0.079728	0.042681	0.006033	0.040674	0.084897	-0.028810	0.009788	0.008208	-0.007117	0.022046	-0.024549	0.001291	0.005651	0.004956	0.075024
ERC20 uniq rec token name	-0.018047	0.047202	-0.014966	0.324288	0.220614	0.202747	0.030089	0.148851	0.237520	-0.002695	0.179039	0.200168	-0.004092	0.165062	0.045800	0.019074	0.012411	0.016543	0.267906

```
In [13]: #Using seaborn
plt.figure(figsize = (30,30))
sns.heatmap(df.corr(), annot = True)
```

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f80640f4cd0>



CHECKING IF DATASET IS SKEWED OR NOT


Histogram

1. A histogram is an approximate representation of the distribution of numerical data.
2. To construct a histogram, the first step is to "bin" (or "bucket") the range of values—that is, divide the entire range of values into a series of intervals—and then count how many values fall into each interval.
3. The words used to describe the patterns in a histogram are: "symmetric", "skewed left" or "right", "unimodal", "bimodal" or "multimodal".

```
In [14]: df.columns[:6]

Out[14]: Index(['FLAG', 'Avg min between sent txn', 'Avg min between received txn',
               'Time Diff between first and last (Mins)', 'Sent txn', 'Received Txn'],
              dtype='object')
```

```
In [15]: # Histogram using pandas
fig, ax = plt.subplots(ncols=4, nrows=1, figsize=(20,5))
index = 0
ax = ax.flatten()
for col, value in df.iloc[:, : 4].items():
    if df[col].dtypes != 'O':
        sns.histplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



Skewness

- The skew method returns a scalar value representing the skewness of the distribution. A positive value indicates a positive skew (i.e., the tail on the right side of the distribution is longer), a negative value indicates a negative skew (i.e., the tail on the left side of the distribution is longer), and a value of 0 indicates that the distribution is symmetrical

```
In [16]: # check skewness of dataset
df.skew(axis = 0, skipna = True)
```

<ipython-input-16-77a6948006c4>:2: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.skew(axis = 0, skipna = True)
```

```
Out[16]: FLAG 1.342100
Avg min between sent txn 8.420000
Avg min between received txn 6.745298
Time Diff between first and last (Mins) 1.809977
Sent txn 10.484544
Received Txn 8.821383
Number of Created Contracts 51.720220
Unique Received From Addresses 18.116107
Unique Sent To Addresses 18.354325
min value received 23.295883
max value received 46.423682
avg val received 96.513680
min val sent 73.426748
max val sent 59.833713
avg val sent 25.531087
min value sent to contract 79.849427
max val sent to contract 78.848894
avg value sent to contract 63.607723
total transactions (including txn to create contract 6.849046
total ether sent 62.363237
total ether received 58.795284
total ether sent contracts 78.848737
total ether balance -1.205262
Total ERC20 txns 19.930233
ERC20 total ether received 94.806604
ERC20 total ether sent 94.786924
ERC20 total ether sent contract 59.280265
ERC20 uniq sent addr 40.648025
ERC20 uniq rec addr 37.586021
ERC20 uniq sent addr.1 23.720828
ERC20 uniq rec contract addr 16.330917
ERC20 avg time between sent txn 0.000000
ERC20 avg time between rec txn 0.000000
ERC20 avg time between rec 2 txn 0.000000
ERC20 avg time between contract txn 0.000000
ERC20 min val rec 50.535806
ERC20 max val rec 94.833841
ERC20 avg val rec 71.026899
ERC20 min val sent 94.882248
ERC20 max val sent 94.904137
ERC20 avg val sent 94.919941
ERC20 min val sent contract 0.000000
ERC20 max val sent contract 0.000000
ERC20 avg val sent contract 0.000000
ERC20 uniq sent token name 12.305550
ERC20 uniq rec token name 15.538524
dtype: float64
```

After checking skewness of dataset and from histogram we can see that the dataset is heavily skewed, with most of the weight being on the left tail.


CHECKING IF DATASET IS BALANCED OR NOT

```
In [17]: # count plot of flag column
sns.countplot(df['FLAG'])
```

/usr/local/lib/python3.8/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f805b270580>
```



```
In [18]: # check total number of fraudulent and non-fraudulent instances
df['FLAG'].value_counts()
```

```
Out[18]: 0    7662
         1    2179
         Name: FLAG, dtype: int64
```

```
In [19]: # find Percentage of non-fraudulent instances
print(f"% of fraudulent instances : {round(df['FLAG'].value_counts(normalize = True)[0]*100,0)}%")

% of fraudulent instances : 78.0%
```

```
In [20]: # find Percentage of fraudulent instances
print(f"% of fraudulent instances : {round(df['FLAG'].value_counts(normalize = True)[1]*100,0)}%")

% of fraudulent instances : 22.0%
```

- The dataset is heavily imbalanced with only 22% of fraudulent instance.

Understanding more about dataset

```
In [21]: # display unique values in categorical columns
for col in df.columns:
    if df[col].dtypes == 'O':
        print(df[col].unique)
```

```
<bound method Series.unique of 0          Cofoundit
1          Livepeer Token
2              None
3          Raiden
4      StatusNetwork
...
9836
9837
9838
9839          NaN
9840
Name: ERC20 most sent token type, Length: 9841, dtype: object>
<bound method Series.unique of 0          Numeraire
1          Livepeer Token
2              XENON
3              XENON
4              EOS
...
9836          GSENetwork
9837      Blockwell say NOTSAFU
9838      Free BOB Tokens - BobsRepair.com
9839          NaN
9840          INS Promo1
Name: ERC20_most_rec_token_type, Length: 9841, dtype: object>
```

```
In [22]: # drop categorical columns
for col in df.columns:
    if df[col].dtypes == 'O':
        df.drop(columns = [col], inplace = True)
```

- Most of the tokens occur only once so they are irrelevant in fraud detection. So dropped them.

```
In [23]: # Replace missings of numerical variables with median
for col in df.columns:
    df[col] = df[col].fillna(df[col].median())
```

```
In [24]: # Filtering the features with 0 variance
variance = df.var()

# Select only the features with a non-zero variance
df_filtered = df[variance[variance > 0].index]
# Drop features with 0 variance --- these features will not help in the performance of the model
df = df.drop(columns = df[variance[variance == 0].index].columns, axis =1)
```

```
In [25]: df_copied = df.copy()
```

```
In [26]: # drop columns that holds only zeros and highly correlated features
df.drop(columns = ['min value sent to contract', 'max val sent to contract','avg value sent to contract', ' ERC20 uniq sent addr.1'], axis = 1, inplace = True)
upper_tri = df.corr().where(np.triu(np.ones(df.corr().shape),k=1).astype(np.bool))
to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.5)]
df = df.drop(columns = to_drop, axis=1)

<ipython-input-26-967f9e24f774>:3: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not
modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
upper_tri = df.corr().where(np.triu(np.ones(df.corr().shape),k=1).astype(np.bool))
```

```
In [27]: # Check which columns are having categorical, numerical or boolean values  ERC20 uniq sent addr.1 'avg value sent to contract'
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9841 entries, 0 to 9840
Data columns (total 20 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   FLAG                                     9841 non-null   int64
1   Avg min between sent tnx                9841 non-null   float64
2   Avg min between received tnx            9841 non-null   float64
3   Time Diff between first and last (Mins) 9841 non-null   float64
4   Sent tnx                                9841 non-null   int64
5   Received Tnx                             9841 non-null   int64
6   Number of Created Contracts              9841 non-null   int64
7   min value received                      9841 non-null   float64
8   max value received                      9841 non-null   float64
9   min val sent                            9841 non-null   float64
10  max val sent                             9841 non-null   float64
11  total Ether sent                        9841 non-null   float64
12  total ether sent contracts               9841 non-null   float64
13  total ether balance                     9841 non-null   float64
14  Total ERC20 txns                        9841 non-null   float64
15  ERC20 total Ether received              9841 non-null   float64
16  ERC20 total ether sent                  9841 non-null   float64
17  ERC20 total Ether sent contract         9841 non-null   float64
18  ERC20 uniq rec contract addr            9841 non-null   float64
19  ERC20 min val rec                       9841 non-null   float64
dtypes: float64(16), int64(4)
memory usage: 1.5 MB
```

In [28]: `#recheck the correlation matrix`
`df.corr()`

Out[28]:

	FLAG	Avg min between sent txn	Avg min between received txn	Time Diff between first and last (Mins)	Sent txn	Received Txn	Number of Created Contracts	min value received	max value received	min val sent	max val sent	total Ether sent	total ether sent contracts	total ether balance	Total ERC20 tnxs	ERC20 total Ether received	ERC20 total ether sent	ERC20 total Ether sent contract	ERC20 uniq rec contract addr	ERC20 min v rec
FLAG	1.000000	-0.029754	-0.118533	-0.269354	-0.078006	-0.079316	-0.013711	-0.021641	-0.019259	0.006626	-0.022437	-0.014993	-0.007988	-0.003229	-0.034132	-0.005711	0.018428	0.008127	-0.052473	0.004432
Avg min between sent txn	-0.029754	1.000000	0.060979	0.214722	-0.032289	-0.035735	-0.006186	-0.014886	-0.007104	-0.000789	-0.005716	-0.006440	0.001044	-0.001425	-0.011061	-0.002339	-0.001916	-0.004101	0.047946	0.004982
Avg min between received txn	-0.118533	0.060979	1.000000	0.303897	-0.040419	-0.053478	-0.008378	-0.045753	-0.011575	-0.008761	-0.012176	-0.009442	-0.003230	-0.002149	-0.019177	-0.003430	-0.003653	-0.005813	-0.011693	-0.007775
Time Diff between first and last (Mins)	-0.269354	0.214722	0.303897	1.000000	0.154480	0.148376	-0.003881	-0.084996	-0.002240	-0.013107	0.014194	0.012999	0.036162	0.002955	0.078482	0.046570	-0.004338	-0.007337	0.324088	-0.008922
Sent txn	-0.078006	-0.032289	-0.040419	0.154480	1.000000	0.198455	0.320603	0.024015	0.102109	-0.004846	0.225356	0.244434	0.013263	-0.127211	0.381859	0.006840	-0.000158	-0.002652	0.221971	-0.003482
Received Txn	-0.079316	-0.035735	-0.053478	0.148376	0.198455	1.000000	-0.003838	-0.022936	0.224805	0.093448	0.097769	0.132150	0.008243	0.158146	0.117093	0.016648	-0.000283	0.104939	0.204128	-0.004042
Number of Created Contracts	-0.013711	-0.006186	-0.008378	-0.003881	0.320603	-0.003838	1.000000	-0.001542	-0.000752	-0.000913	0.141919	0.024877	-0.000225	-0.037902	0.249601	0.002099	0.001192	-0.000334	0.029421	-0.000722
min value received	-0.021641	-0.014886	-0.045753	-0.084996	0.024015	-0.022936	-0.001542	1.000000	0.026710	0.117682	0.015061	-0.001748	-0.002015	-0.000233	-0.006751	-0.001580	-0.001488	-0.002331	-0.000881	-0.000842
max value received	-0.019259	-0.007104	-0.011575	-0.002240	0.102109	0.224805	-0.000752	0.026710	1.000000	0.002239	0.135937	0.112739	-0.000550	0.278315	0.032524	0.006272	-0.000225	0.068761	0.181755	-0.000972
min val sent	0.006626	-0.000789	-0.008761	-0.013107	-0.004846	0.093448	-0.000913	0.117682	0.002239	1.000000	0.022662	0.000166	-0.000519	0.001514	-0.002451	-0.000406	-0.000389	-0.000600	-0.003749	-0.000442
max val sent	-0.022437	-0.005716	-0.012176	0.014194	0.225356	0.097769	0.141919	0.015061	0.135937	0.022662	1.000000	0.470139	-0.000594	-0.564872	0.131102	0.032925	0.000663	-0.000822	0.166292	-0.001002
total Ether sent	-0.014993	-0.006440	-0.009442	0.012999	0.244434	0.132150	0.024877	-0.001748	0.112739	0.000166	0.470139	1.000000	-0.000356	-0.313816	0.072154	0.003749	0.000199	-0.000491	0.088149	-0.000775
total ether sent contracts	-0.007988	0.001044	-0.003230	0.036162	0.013263	0.008243	-0.000225	-0.002015	-0.000550	-0.000519	-0.000594	-0.000356	1.000000	-0.000074	-0.000201	0.000442	-0.000165	-0.000259	0.012325	-0.000442
total ether balance	-0.003229	-0.001425	-0.002149	0.002955	-0.127211	0.158146	-0.037902	-0.000233	0.278315	0.001514	-0.564872	-0.313816	-0.000074	1.000000	-0.007995	-0.001737	-0.000301	0.037944	-0.023787	-0.000172
Total ERC20 tnxs	-0.034132	-0.011061	-0.019177	0.078482	0.381859	0.117093	0.249601	-0.006751	0.032524	-0.002451	0.131102	0.072154	-0.000201	-0.007995	1.000000	0.004378	0.000802	0.006198	0.264907	-0.001982
ERC20 total Ether received	-0.005711	-0.002339	-0.003430	0.046570	0.006840	0.016648	0.002099	-0.001580	0.006272	-0.000406	0.032925	0.003749	0.000442	-0.001737	0.004378	1.000000	0.000113	-0.000132	0.030039	-0.000322
ERC20 total ether sent	0.018428	-0.001916	-0.003653	-0.004338	-0.000158	-0.000283	0.001192	-0.001488	-0.000225	-0.000389	0.000663	0.000199	-0.000165	-0.000301	0.000802	0.000113	1.000000	-0.000195	0.001920	-0.000252
ERC20 total ether sent contract	0.008127	-0.004101	-0.005813	-0.007337	-0.002652	0.104939	-0.000334	-0.002331	0.068761	-0.000600	-0.000822	-0.000491	-0.000259	0.037944	0.006198	-0.000132	-0.000195	1.000000	0.037021	-0.000442
ERC20 uniq rec contract addr	-0.052473	0.047946	-0.011693	0.324088	0.221971	0.204128	0.029421	-0.000881	0.181755	-0.003749	0.166292	0.088149	0.012325	-0.023787	0.264907	0.030039	0.001920	0.037021	1.000000	-0.005930
ERC20 min val rec	0.004434	0.004998	-0.007794	-0.008921	-0.003480	-0.004043	-0.000724	-0.000847	-0.000976	-0.000446	-0.001008	-0.000766	-0.000412	-0.000170	-0.001969	-0.000322	-0.000297	-0.000477	-0.005930	1.000000

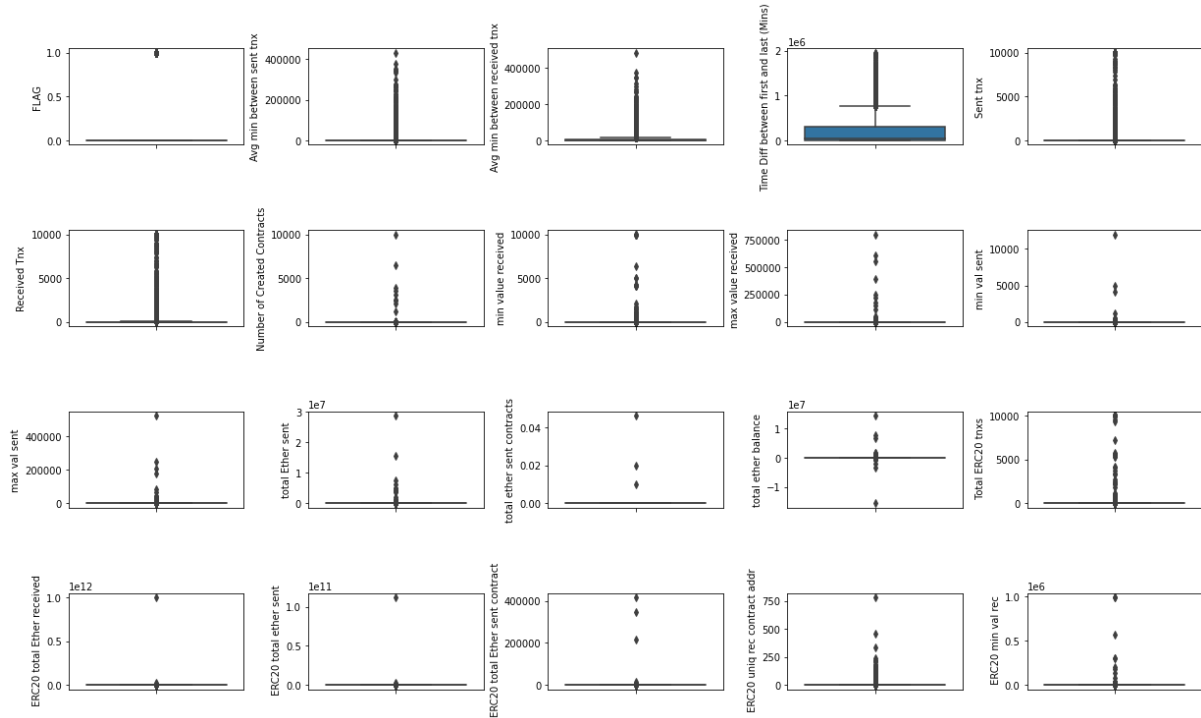
In [29]: `# check columns of dataset`
`df.columns`

Out[29]: Index(['FLAG', 'Avg min between sent txn', 'Avg min between received txn', 'Time Diff between first and last (Mins)', 'Sent txn', 'Received Txn', 'Number of Created Contracts', 'min value received', 'max value received', 'min val sent', 'max val sent', 'total Ether sent', 'total ether sent contracts', 'total ether balance', 'Total ERC20 tnxs', 'ERC20 total Ether received', 'ERC20 total ether sent', 'ERC20 total Ether sent contract', 'ERC20 uniq rec contract addr', 'ERC20 min val rec'], dtype='object')

Box plot

- A boxplot is a standardized way of displaying the dataset based on a five-number summary:
 - Minimum (Q0 or 0th percentile): the lowest data point excluding any outliers.
 - Maximum (Q4 or 100th percentile): the largest data point excluding any outliers.
 - Median (Q2 or 50th percentile): the middle value of the dataset.
 - First quartile (Q1 or 25th percentile): also known as the lower quartile $q_n(0.25)$, is the median of the lower half of the dataset.
 - Third quartile (Q3 or 75th percentile): also known as the upper quartile $q_n(0.75)$, is the median of the upper half of the dataset

```
# box plot using pandas
# box plot for Avg min between sent tnx column
fig, ax = plt.subplots(ncols = 5, nrows = 4, figsize=(17,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    sns.boxplot(y = col, data = df, ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




- From all boxen plot we can see that the dataset is heavily skewed

DENDROGRAM

```
# Plot a Dendrogram on the columns of the dataset
# from collections import defaultdict
from scipy.spatial.distance import pdist, squareform
from scipy.cluster.hierarchy import linkage, dendrogram

%pylab inline
from pylab import rcParams
rcParams['figure.figsize'] = 12, 9
clustdf_t = df.transpose()
c_dist = pdist(clustdf_t) # computing the distance
c_link = linkage(clustdf_t, metric='correlation', method='complete') # computing the Linkage
B=dendrogram(c_link, labels=list(df.columns),orientation='left')
# dropping the NaN values
```

Populating the interactive namespace from numpy and matplotlib



DATA PREPROCESSING

```
In [32]: ## Split the labels and the target
X = df.iloc[ : , 1: ]
y = df.iloc[ : , :1 ]

#check the shape
print(X.shape)
print(y.shape)

(9841, 19)
(9841, 1)

In [33]: # import train test split
from sklearn.model_selection import train_test_split

In [34]: # Split into training (80%) and testing set (20%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
# check the shape of train and test
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(7872, 19)
(1969, 19)
(7872, 1)
(1969, 1)
```

HANDLING SKEWNESS

```
In [35]: # import pipeline
#import standar scalar
# import simple imputer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer

In [36]: # create pipeline
# 1.simple imputer
# 2. standard scalar
pipe = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])

In [37]: # fit the pipeline on train data
X_train = pipe.fit_transform(X_train)
X_test = pipe.transform(X_test)
```

HANDLING IMBALANCE

SMOTE stands for Synthetic Minority Oversampling Technique. This is a statistical technique for increasing the number of cases in your dataset in a balanced way. The module works by generating new instances from existing minority cases that you supply as input.

SMOTE - <https://www.geeksforgeeks.org/ml-handling-imbalanced-data-with-smote-and-near-miss-algorithm-in-python/> (<https://www.geeksforgeeks.org/ml-handling-imbalanced-data-with-smote-and-near-miss-algorithm-in-python/>)

```
In [38]: from imblearn.over_sampling import SMOTE

# Instantiate the SMOTE class
sm = SMOTE(sampling_strategy='minority')

# Fit and transform the training data using SMOTE
X_train, y_train = sm.fit_resample(X_train, y_train)

# Print the shape of X_train after oversampling
print("Shape of X_train after oversampling:", X_train.shape)

# Print the shape of y_train after oversampling
print("Shape of y_train after oversampling:", y_train.shape)

Shape of X_train after oversampling: (12222, 19)
Shape of y_train after oversampling: (12222, 1)
```

MODELING

```
In [39]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from imblearn.under_sampling import RandomUnderSampler
from imblearn.pipeline import Pipeline
from sklearn.linear_model import SGDClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
# Evaluation Metrics
from sklearn.metrics import accuracy_score, roc_auc_score, classification_report, precision_score, recall_score, f1_score, confusion_matrix, roc_curve
# Plots
from sklearn import tree
```



```
In [40]: # Function for calculating all the relevant metrics
        """Function to calculate all evaluation metrics"""
        def evaluation(y_test,y_pred):
            Accuracy = accuracy_score(y_test, y_pred)
            conf_mat = confusion_matrix(y_test, y_pred)
            true_positive = conf_mat[0][0]
            false_positive = conf_mat[0][1]
            false_negative = conf_mat[1][0]
            true_negative = conf_mat[1][1]
            Precision = true_positive / (true_positive + false_positive)
            Recall = true_positive / (true_positive + false_negative)
            F1_Score = 2 * (Recall * Precision) / (Recall + Precision)
            AUC = roc_auc_score(y_test, y_pred)
            return Accuracy, Precision, Recall, F1_Score, AUC
```

Model Creation

Hyper parameter tuning

- A hyperparameter is a parameter whose value is set before the learning process begins
- Hyperparameters tuning is crucial as they control the overall behavior of a machine learning model
- Every machine learning models will have different hyperparameters that can be set

Note: Directly hyperparameter tuning has been done to get the best accuracy based on the experience. Beginner should always build basic model first then go for hyper parameter tuning

```
In [41]: '''Hyperparameters of Logistic Regression'''
        LR_parameters = {
            'C': np.random.uniform(0.001,1,5),
            'tol': np.random.uniform(0.0001,0.1,5)
        }

        '''Hyperparameters of SGD Classifier With Hinge Loss'''
        SGD_parameters = {
            'penalty': ['l1', 'l2'],
            'alpha': np.random.uniform(0.0001,0.1,9)
        }

        '''Hyperparameters of Naive Bayes'''
        NB_parameters = {
            'var_smoothing': np.random.uniform(1e-16,1e-14,100)
        }

        '''Hyperparameters of Decision Tree Classifier'''
        DTC_parameters = {
            'max_depth': np.random.randint(10,25,5)
        }

        '''Hyperparameters of Random Forest Classifier'''
        RF_parameters = {
            'n_estimators': [300,500,600,650,700],
            'max_depth': [80,110,125,135]
        }

        '''Hyperparameters of Gradient Boosted Decision Trees'''
        GBDT_parameters = {
            'n_estimators': [100, 250, 350, 500],
            'max_depth': [4, 6, 10, 15],
            'learning_rate': [0.001, 0.01, 0.1, 1, 10]
        }

        '''Hyperparameters of AdaBoost Classifier'''
        AB_parameters = {
            'n_estimators': [ 400 , 600 , 650 , 700 , 750 , 800 ]
        }

        '''ALL Models'''
        models = {
            1: RandomizedSearchCV(estimator = LogisticRegression(), param_distributions = LR_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5),
            2: RandomizedSearchCV(estimator = SGDClassifier(), param_distributions = SGD_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5),
            3: RandomizedSearchCV(estimator = GaussianNB(), param_distributions = NB_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5),
            4: RandomizedSearchCV(estimator = DecisionTreeClassifier(), param_distributions = DTC_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5),
            5: RandomizedSearchCV(estimator = RandomForestClassifier(), param_distributions = RF_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5),
            6: RandomizedSearchCV(estimator = GradientBoostingClassifier(), param_distributions = GBDT_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5),
            7: RandomizedSearchCV(estimator = AdaBoostClassifier(), param_distributions = AB_parameters, scoring= 'f1', verbose=2, n_jobs = -1, cv=5)
        }

        map_keys = list(models.keys())
```

```
In [42]: # Get model name using id from linear_model_collection
        def get_model_building_technique_name(num):
            if num == 1:
                return 'LogisticRegression_Tuned'
            if num == 2:
                return 'SGDClassifier_Hinge_Loss_Tuned'
            if num == 3:
                return 'NaiveBayes_Tune'
            if num == 4:
                return 'DecisionTreeClassifier_Tuned'
            if num == 5:
                return 'RandomForestClassifier_Tuned'
            if num == 6:
                return 'GradientBoostingClassifier_Tuned'
            if num == 7:
                return 'AdaBoostClassifier_Tuned'
            return ''
```

```
In [43]: results = []
for key_index in range(len(map_keys)):
    key = map_keys[key_index]
    try:
        # if key in [3,4,5,6,7,8]:
        model = models[key]
        print(key)
        model.fit(X_train, y_train)

        y_pred_train = model.predict(X_train)
        y_pred_test = model.predict(X_test)

        '''Test Accuracy'''
        Accuracy_Test, Precision_Test, Recall_Test, F1_Score_Test, AUC_Test = evaluation(y_test, y_pred_test)

        '''Train Accuracy'''
        y_pred_train = model.predict(X_train)

        Accuracy_Train, Precision_Train, Recall_Train, F1_Score_Train, AUC_Train = evaluation(y_train, y_pred_train)

        results.append({
            'Model Name' : get_model_building_technique_name(key),
            'Trained Model' : model,
            'Accuracy_Test' : Accuracy_Test,
            'Precision_Test' : Precision_Test,
            'Recall_Test' : Recall_Test,
            'F1_Macro_Score_Test' : F1_Score_Test,
            'AUC_Test' : AUC_Test,
            'Accuracy_Train' : Accuracy_Train,
            'Precision_Train' : Precision_Train,
            'Recall_Train' : Recall_Train,
            'F1_Macro_Score_Train' : F1_Score_Train,
            'AUC_Train' : AUC_Train
        })

    except Exception as e:
        print(e)
```

1
Fitting 5 folds for each of 10 candidates, totalling 50 fits

/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

2
Fitting 5 folds for each of 10 candidates, totalling 50 fits

/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

3
Fitting 5 folds for each of 10 candidates, totalling 50 fits

/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/_search.py:292: UserWarning: The total space of parameters 5 is smaller than n_iter=10. Running 5 iterations. For exhaustive searches, use GridSearchCV.
warnings.warn()

4
Fitting 5 folds for each of 5 candidates, totalling 25 fits

5
Fitting 5 folds for each of 10 candidates, totalling 50 fits

/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/_search.py:926: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
self.best_estimator_.fit(X, y, **fit_params)

6
Fitting 5 folds for each of 10 candidates, totalling 50 fits

/usr/local/lib/python3.8/dist-packages/sklearn/ensemble/_gb.py:494: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

7
Fitting 5 folds for each of 6 candidates, totalling 30 fits

/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/_search.py:292: UserWarning: The total space of parameters 6 is smaller than n_iter=10. Running 6 iterations. For exhaustive searches, use GridSearchCV.
warnings.warn()
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

```
In [44]: result_df = pd.DataFrame(results)
result_df
# result_df_test = result_df.iloc[:, [0,2,3,4,5,6]]
# result_df_test
# result_df_train = result_df.iloc[:, [0,7,8,9,10,11]]
# result_df_train
```

Out[44]:

	Model Name	Trained Model	Accuracy_Test	Precision_Test	Recall_Test	F1_Macro_Score_Test	AUC_Test	Accuracy_Train	Precision_Train	Recall_Train	F1_Macro_Score_Train	AUC_Train
0	LogisticRegression_Tuned	RandomizedSearchCV(cv=5, estimator=LogisticReg...	0.604368	0.535783	0.933708	0.680868	0.697318	0.707822	0.552610	0.801376	0.654140	0.707822
1	SGDClassifier_Hinge_Loss_Tuned	RandomizedSearchCV(cv=5, estimator=SGDClassifi...	0.601828	0.534494	0.930415	0.678952	0.693084	0.705367	0.552774	0.795572	0.652312	0.705367
2	NaiveBayes_Tune	RandomizedSearchCV(cv=5, estimator=GaussianNB(...	0.311326	0.128304	0.980296	0.226910	0.559367	0.566519	0.150057	0.898139	0.257151	0.566519
3	DecisionTreeClassifier_Tuned	RandomizedSearchCV(cv=5, estimator=DecisionTre...	0.960386	0.974855	0.974855	0.974855	0.940777	1.000000	1.000000	1.000000	1.000000	1.000000
4	RandomForestClassifier_Tuned	RandomizedSearchCV(cv=5, estimator=RandomFores...	0.974606	0.984526	0.983258	0.983892	0.961163	1.000000	1.000000	1.000000	1.000000	1.000000
5	GradientBoostingClassifier_Tuned	RandomizedSearchCV(cv=5, estimator=GradientBoo...	0.974099	0.982592	0.984496	0.983543	0.962588	1.000000	1.000000	1.000000	1.000000	1.000000
6	AdaBoostClassifier_Tuned	RandomizedSearchCV(cv=5, estimator=AdaBoostCla...	0.965465	0.970342	0.985593	0.977908	0.958855	0.984045	0.982163	0.985874	0.984015	0.984045

ROC Curve

- The overall performance of a classifier, summarized over all possible thresholds, is given by the Receiver Operating Characteristics (ROC) curve. The name "ROC" is historical and comes from communications theory. ROC Curves are used to see how well your classifier can separate positive and negative examples and to identify the best threshold for separating them.
- To be able to use the ROC curve, your classifier should be able to rank examples such that the ones with higher rank are more likely to be positive (fraudulent)
- ROC Curves summarize the trade-off between the true positive rate and false positive rate for a predictive model using different probability thresholds.

AUC (Area Under the Curve)


- The model performance is determined by looking at the area under the ROC curve (or AUC). An excellent model has AUC near to the 1.0, which means it has a good measure of separability

```
In [54]: fpr_dict = {}
         tpr_dict = {}
         for i in range(len(map_keys)):


             model_pred = result_df['Trained Model'][i].predict(X_test)
             fpr, tpr, thresholds = roc_curve(y_test, model_pred)
             fpr_dict[i] = fpr
             tpr_dict[i] = tpr

         plt.figure(figsize=(15,8))
         plt.suptitle('\nReceiver Operating Characteristic (ROC) Curve', fontsize=20)
         plt.plot(fpr_dict[0], tpr_dict[0], color='orange', label=f"ROC - {result_df['Model Name'][0]}")
         plt.plot(fpr_dict[1], tpr_dict[1], color='green', label=f"ROC - {result_df['Model Name'][1]}")
         plt.plot(fpr_dict[2], tpr_dict[2], color='yellow', label=f"ROC - {result_df['Model Name'][2]}")
         plt.plot(fpr_dict[3], tpr_dict[3], color='red', label=f"ROC - {result_df['Model Name'][3]}")
         plt.plot(fpr_dict[4], tpr_dict[4], color='black', label=f"ROC - {result_df['Model Name'][4]}")
         plt.plot(fpr_dict[5], tpr_dict[5], color='blue', label=f"ROC - {result_df['Model Name'][5]}")
         plt.plot(fpr_dict[6], tpr_dict[6], color='blue', label=f"ROC - {result_df['Model Name'][6]}")

         plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve')
         plt.xlabel('False Positive Rate', fontdict={'fontsize': 20})
         plt.ylabel('True Positive Rate', fontdict={'fontsize': 20})
         plt.legend()
         plt.show()
```



```
In [55]: plt.figure(figsize=(15,8))
plt.suptitle('\nRoc-Auc-Score Distribution\n\n', fontsize=4, fontweight='bold')
sns.barplot(data=result_df, x='Model Name', y='AUC_Test')
plt.xlabel('Model Name',fontdict={'fontsize': 20})
plt.ylabel('Roc-Auc-Score',fontdict={'fontsize': 20})
plt.xticks(rotation=45)
plt.show()
```



```
In [56]: Best_Model_Name = result_df['Trained Model'][result_df[result_df['AUC_Test'] == max(result_df['AUC_Test'])]['Trained Model'].index[0]
Best_Model_Index = result_df['Trained Model'][result_df[result_df['AUC_Test'] == max(result_df['AUC_Test'])]['Trained Model'].index].index[0]
Best_Model_Name
```

```
Out[56]: RandomizedSearchCV(cv=5, estimator=GradientBoostingClassifier(), n_jobs=-1,
param_distributions={'learning_rate': [0.001, 0.01, 0.1, 1,
10],
'max_depth': [4, 6, 10, 15],
'n_estimators': [100, 250, 350, 500]},
scoring='f1', verbose=2)
```

```
In [57]: import pickle
with open('ETHEREUM_Fraud_Detection.sav', 'wb') as best_model_pickle:
pickle.dump(Best_Model_Name, best_model_pickle)
```

Conclusion

- We did training and prediction using all the above models and selected Gradient Boosting Classifier as final model as it performed well compared to other models with accuracy of 98%
- We have performed EDA, preprocessing, build different models, visualized feature importance, hyper parameter tuning and did prediction
- We also perform necessary operations to handle imbalanced and skewed nature of data

THE END